

## Regression Assignment: v1

**Steps:** Import all the necessary Libraries

\* First 5 Rows of the dataset:

S.no	age	sex	bmi	children	smoker	charges
0	19	female	27.9	0	yes	16884.92
1	18	male	33.77	1	no	1725.552
2	28	male	33	3	no	4449.462
3	33	male	22.705	0	no	21984.47
4	32	male	28.88	0	no	3866.855

### 1. Identifying the problem statement:

- Stage 1:

\* We have Input data's are Numbers - so Domain will be Machine Learning

- Stage 2:

\* Supervised Learning- It's a Clear requirement and we have IP/OP data.

- stage 3:

\* Numerical values (Target features) - so we can go with Regression

### 2. Basic info about the dataSet:

\* Total Row : 1388

\* Total Columns : 6

\* Dataset contains both numerical and categorical values

\* Column names : AGE, SEX, BMI, CHILDREN, SMOKER, CHARGES(Target variable)

\* age,children are Integer type.

\* bmi and charges are Float type.

\* sex and smoker are in Object type.

\* There is no null values in the dataset.

\* Requirement is clear.

### 3.Preprocessing Method:

Converting categorical variable to continuous numerical variable

After Converting data:

s.no	age	bmi	children	charges	sex_male	smoker_yes
0	19	27.9	0	16884.924	0	1

1	18	33.77	1	1725.5523	1	0
2	28	33	3	4449.462	1	0
3	33	22.705	0	21984.47061	1	0
4	32	28.88	0	3866.8552	1	0

- Separate the data as independent and dependent
- Split the data using train\_test\_split concept
- If the data is looks too different within the independent variable then go for Scaling concept

**4. Create different Models to predict the charges and choose the best model with high r2\_score value.**

a) Multiple Linear Regression:

Results: **R2\_Score = 0.77239**


b) SVM Regression: Default value:

## Model 2 - SVR

[131]:

```
from sklearn.svm import SVR
svr_model = SVR()
svr_model.fit(X_train, y_train)
```

[131]:



[133]:

```
y_pred = svr_model.predict(X_test)
```

[135]:

```
svr_r2score = r2_score(y_test, y_pred)
svr_r2score
```

[135]:

-0.09743369320622963

Model with Different parameters

**Best : RBF, C=3000 and R2\_Score = 0.84223**

[187]:

```
# SVR_R2_Score : Results with all the combinations:
for kernel, C, score in svr_r2score:
    print("Kernel:",kernel,"- C:",C,"- R² Score",score)
```

```
Kernel: linear - C: 10 - R² Score 0.4411154085859137
Kernel: linear - C: 100 - R² Score 0.6205250918367728
Kernel: linear - C: 500 - R² Score 0.7486844603269274
Kernel: linear - C: 1000 - R² Score 0.7187303718390283
Kernel: linear - C: 2000 - R² Score 0.717185774596472
Kernel: linear - C: 3000 - R² Score 0.7161570985448058
Kernel: rbf - C: 10 - R² Score -0.044010385361080706
Kernel: rbf - C: 100 - R² Score 0.3229298217686596
Kernel: rbf - C: 500 - R² Score 0.6597447201916936
Kernel: rbf - C: 1000 - R² Score 0.7923789012887561
Kernel: rbf - C: 2000 - R² Score 0.8332041538541531
Kernel: rbf - C: 3000 - R² Score 0.8422384746495604
Kernel: poly - C: 10 - R² Score 0.024580030588300605
Kernel: poly - C: 100 - R² Score 0.5897995020893756
Kernel: poly - C: 500 - R² Score 0.8038815140652048
Kernel: poly - C: 1000 - R² Score 0.8338303273683121
Kernel: poly - C: 2000 - R² Score 0.8403113295667679
Kernel: poly - C: 3000 - R² Score 0.8390286005966449
Kernel: sigmoid - C: 10 - R² Score 0.022978265830825184
Kernel: sigmoid - C: 100 - R² Score 0.5099640007962808
Kernel: sigmoid - C: 500 - R² Score 0.46829640989809973
Kernel: sigmoid - C: 1000 - R² Score 0.3012743569891182
Kernel: sigmoid - C: 2000 - R² Score -0.8798843869243349
Kernel: sigmoid - C: 3000 - R² Score -2.9047335174024607
```

Activate Windows

### c) Decision Tree Regressor:

With Default Values score:

## Model 3 - Decision Tree Regressor

[187]:

```
from sklearn.tree import DecisionTreeRegressor
DT_model = DecisionTreeRegressor()
DT_model.fit(X_train,y_train)
y_pred = DT_model.predict(X_test)
DT_r2 = r2_score(y_test,y_pred)
DT_r2
```

[187]:

0.7064796314633288

With different parameters:

## Model 3 - Decision Tree Regressor

```
: from sklearn.tree import DecisionTreeRegressor
criterion1 = ['squared_error', 'friedman_mse', 'absolute_error', 'poisson']
max_ft1 = ['sqrt', 'log2']
splitter1 = ['best', 'random']
DT_r2 = []

for cri in criterion1:
    for mx_ft in max_ft1:
        for split in splitter1:
            DT_model = DecisionTreeRegressor(criterion=cri, max_features=mx_ft, splitter=split)
            DT_model.fit(X_train, y_train)
            y_pred = DT_model.predict(X_test)
            dt_r2 = r2_score(y_test, y_pred)
            DT_r2.append((cri, mx_ft, split, dt_r2))
```

```
for criterion, max_features, splitter, score in DT_r2:
    print("Criterion:", criterion, "max_features:", max_features, "splitter:", splitter, "R2_Score:", score)
```

```
Criterion: squared_error max_features: sqrt splitter: best R2_Score: 0.6905149345691342
Criterion: squared_error max_features: sqrt splitter: random R2_Score: 0.6089158084430758
Criterion: squared_error max_features: log2 splitter: best R2_Score: 0.7208084204160041
Criterion: squared_error max_features: log2 splitter: random R2_Score: 0.5885137699183175
Criterion: friedman_mse max_features: sqrt splitter: best R2_Score: 0.741628247754792
Criterion: friedman_mse max_features: sqrt splitter: random R2_Score: 0.632694706926831
Criterion: friedman_mse max_features: log2 splitter: best R2_Score: 0.6353895211454534
Criterion: friedman_mse max_features: log2 splitter: random R2_Score: 0.6909830368433382
Criterion: absolute_error max_features: sqrt splitter: best R2_Score: 0.6465435342030181
Criterion: absolute_error max_features: sqrt splitter: random R2_Score: 0.7744421292327105
Criterion: absolute_error max_features: log2 splitter: best R2_Score: 0.712613522330376
Criterion: absolute_error max_features: log2 splitter: random R2_Score: 0.6719344044572477
Criterion: poisson max_features: sqrt splitter: best R2_Score: 0.5326888585822843
Criterion: poisson max_features: sqrt splitter: random R2_Score: 0.6616041177188116
Criterion: poisson max_features: log2 splitter: best R2_Score: 0.5525323495374564
Criterion: poisson max_features: log2 splitter: random R2_Score: 0.6235679214227591
```

After Using diff parameters the best **R2\_Score** is : **0.774421**

D) Random Forest Regressor: with Default values

## Random\_Forest\_Regressor

[204]:

```
from sklearn.ensemble import RandomForestRegressor
rff = RandomForestRegressor()
rff.fit(X_train, y_train)
y_pred = rff.predict(X_test)
randomforest_r2score = r2_score(y_test, y_pred)
randomforest_r2score
```

[204]:

0.8373310073237741

1 1 1

With different parameters:

## Random\_Forest\_Regressor

```
08]: from sklearn.ensemble import RandomForestRegressor

n_estimators = [50,100,200,250,300,350,400,500]
criterion2 = ['squared_error','friedman_mse','poisson']
max_ft2 = ['sqrt','log2']
randomforest_r2score= []

for estimators in n_estimators:
    for crt in criterion2:
        for max_frt in max_ft2:
            rff = RandomForestRegressor(n_estimators=estimators ,criterion=crt ,max_features=max_frt ,random_state=32)
            rff.fit(X_train,y_train)
            y_pred = rff.predict(X_test)
            rf_r2score = r2_score(y_test,y_pred)
            randomforest_r2score.append((estimators,crt,max_frt,rf_r2score))

09]: for n_estimators,criterion,max_features,score in randomforest_r2score:
    print("n_estimators:",n_estimators,"Criterion:",criterion,"max_features:",max_features,"R2_Score:",score)
```

n\_estimators: 50 Criterion: squared\_error max\_features: sqrt R2\_Score: 0.8502984279389878  
n\_estimators: 50 Criterion: squared\_error max\_features: log2 R2\_Score: 0.8502984279389878  
n\_estimators: 50 Criterion: friedman\_mse max\_features: sqrt R2\_Score: 0.8509127600875582  
n\_estimators: 50 Criterion: friedman\_mse max\_features: log2 R2\_Score: 0.8509127600875582  
n\_estimators: 50 Criterion: poisson max\_features: sqrt R2\_Score: 0.8468833336548729  
n\_estimators: 50 Criterion: poisson max\_features: log2 R2\_Score: 0.8468833336548729  
n\_estimators: 100 Criterion: squared\_error max\_features: sqrt R2\_Score: 0.8492395564499937  
n\_estimators: 100 Criterion: squared\_error max\_features: log2 R2\_Score: 0.8492395564499937  
n\_estimators: 100 Criterion: friedman\_mse max\_features: sqrt R2\_Score: 0.8492782583516395  
n\_estimators: 100 Criterion: friedman\_mse max\_features: log2 R2\_Score: 0.8492782583516395  
n\_estimators: 100 Criterion: poisson max\_features: sqrt R2\_Score: 0.8474801406941771  
n\_estimators: 100 Criterion: poisson max\_features: log2 R2\_Score: 0.8474801406941771  
n\_estimators: 200 Criterion: squared\_error max\_features: sqrt R2\_Score: 0.85125883509619  
n\_estimators: 200 Criterion: squared\_error max\_features: log2 R2\_Score: 0.85125883509619  
n\_estimators: 200 Criterion: friedman\_mse max\_features: sqrt R2\_Score: 0.8513141835838371  
n\_estimators: 200 Criterion: friedman\_mse max\_features: log2 R2\_Score: 0.8513141835838371  
n\_estimators: 200 Criterion: poisson max\_features: sqrt R2\_Score: 0.849928342290375  
n\_estimators: 200 Criterion: poisson max\_features: log2 R2\_Score: 0.849928342290375  
n\_estimators: 250 Criterion: squared\_error max\_features: sqrt R2\_Score: 0.8509722090294829  
n\_estimators: 250 Criterion: squared\_error max\_features: log2 R2\_Score: 0.8509722090294829  
n\_estimators: 250 Criterion: friedman\_mse max\_features: sqrt R2\_Score: 0.8510874568342656  
n\_estimators: 250 Criterion: friedman\_mse max\_features: log2 R2\_Score: 0.8510874568342656  
n\_estimators: 250 Criterion: poisson max\_features: sqrt R2\_Score: 0.8500928704645765  
n\_estimators: 250 Criterion: poisson max\_features: log2 R2\_Score: 0.8500928704645765  
n\_estimators: 300 Criterion: squared\_error max\_features: sqrt R2\_Score: 0.8515867676508716  
n\_estimators: 300 Criterion: squared\_error max\_features: log2 R2\_Score: 0.8515867676508716  
n\_estimators: 300 Criterion: friedman\_mse max\_features: sqrt R2\_Score: 0.8516592052895238  
n\_estimators: 300 Criterion: friedman\_mse max\_features: log2 R2\_Score: 0.8516592052895238  
n\_estimators: 300 Criterion: poisson max\_features: sqrt R2\_Score: 0.8509948045947353  
n\_estimators: 300 Criterion: poisson max\_features: log2 R2\_Score: 0.8509948045947353  
n\_estimators: 350 Criterion: squared\_error max\_features: sqrt R2\_Score: 0.8510703568787773  
n\_estimators: 350 Criterion: squared\_error max\_features: log2 R2\_Score: 0.8510703568787773  
n\_estimators: 350 Criterion: friedman\_mse max\_features: sqrt R2\_Score: 0.851144935393941  
n\_estimators: 350 Criterion: friedman\_mse max\_features: log2 R2\_Score: 0.851144935393941  
n\_estimators: 350 Criterion: poisson max\_features: sqrt R2\_Score: 0.8502215755711253  
n\_estimators: 350 Criterion: poisson max\_features: log2 R2\_Score: 0.8502215755711253  
n\_estimators: 400 Criterion: squared\_error max\_features: sqrt R2\_Score: 0.8509950319542703  
n\_estimators: 400 Criterion: squared\_error max\_features: log2 R2\_Score: 0.8509950319542703  
n\_estimators: 400 Criterion: friedman\_mse max\_features: sqrt R2\_Score: 0.8512436813151877  
n\_estimators: 400 Criterion: friedman\_mse max\_features: log2 R2\_Score: 0.8512436813151877  
n\_estimators: 400 Criterion: poisson max\_features: sqrt R2\_Score: 0.8508266616553146  
n\_estimators: 400 Criterion: poisson max\_features: log2 R2\_Score: 0.8508266616553146  
n\_estimators: 500 Criterion: squared\_error max\_features: sqrt R2\_Score: 0.8515876208578533

n\_estimators: 500 Criterion: squared\_error max\_features: log2 R2\_Score: 0.8515876208578533  
n\_estimators: 500 Criterion: friedman\_mse max\_features: sqrt R2\_Score: 0.8515351540982423  
n\_estimators: 500 Criterion: friedman\_mse max\_features: log2 R2\_Score: 0.8515351540982423  
n\_estimators: 500 Criterion: poisson max\_features: sqrt R2\_Score: 0.8518107418752442

**Best Param with best R2\_Score :**

n\_estimators: 500  
Criterion: poisson  
max\_features: log2

**R2\_Score: 0.8518107418752442**

**Conclusion:**

- Based on the **R<sup>2</sup> Score** values from the table, I conclude that the **Random Forest Regressor** is the best-performing model.
- Other models (MLR, SVM, Decision Tree) have lower R<sup>2</sup> scores compared to Random Forest model.

sl.no	Model_Name	R2_Score
1	MLR	0.77239
2	SVM_Regressor	0.84223
3	Decision Tree	0.77444
4	Random Forest	0.85181